

ENTERPRISE ARCHITECTURE FRAMEWORK | WP-2026-02

Enterprise AI Data Readiness Framework: Transforming Fragmented Schemas into AI-Ready Gold Standards

Author: KaoinAI Data Engineering Practice & Center for SME Data Modernization
Core Domains: Master Data Management (MDM), Schema Reconciliation, Entity Resolution

Executive Summary: Most small-to-medium enterprises (SMEs) possess substantial data assets distributed across siloed systems: Shopify for e-commerce, Salesforce/HubSpot for CRM, PostgreSQL for core applications, and Zendesk for customer support. When organizations attempt to deploy AI analytics or LLM workflows on top of these un-curated silos, they encounter severe entity collision, orphaned keys, and conflicting revenue definitions. This paper outlines the 5-pillar KaoinAI Data Readiness Model for establishing golden master records in under 48 hours.

1. The 5 Pillars of Enterprise AI Data Readiness

Before deploying LLMs, data assets must satisfy five quantifiable engineering thresholds:

Pillar	Definition	Readiness Metric	Risk If Neglected
1. Semantic Cohesion	Uniform business definitions across tables	> 98% Glossary Match	LLM answers vary depending on which table is queried
2. Entity Unification	Single golden identity across CRM/Billing	< 0.5% Duplicate Rate	Customer churn alerts fire for wrong accounts
3. Schema Determinism	Strict typing, foreign keys, documented constraints	100% Normalized DDL	SQL syntax errors and costly failed query retries
4. Temporal Freshness	Continuous synchronization via CDC or mirroring	Max 15-min Lag	Agents recommend inventory actions on stale stock
5. Privacy Sanitization	Automated redaction of sensitive identifiers	0 Unmasked PII Rows	Statutory regulatory fines up to \$1M under PDPA

2. Autonomous Entity Resolution: The Golden Record Pipeline

SMEs frequently suffer from 'identity fragmentation': the same client appears as 'Acme Pte Ltd' in Salesforce, 'Acme Logistics' in Shopify, and 'darren@acme.com' in Zendesk. KaoinAI implements a dual-stage deterministic and fuzzy Jaro-Winkler entity resolution model. It links multi-source records into an immutable Golden Customer ID with continuous audit trails.

- **Stage 1: Deterministic Exact Keys:** Links phone hashes, national tax IDs (UEN), and exact normalized email addresses.
- **Stage 2: Probabilistic Semantic Matching:** Resolves typographical variances, regional company suffixes ('Pte Ltd' vs 'LLC'), and phonetic transliterations with > 99.2% statistical confidence.
- **Stage 3: Survivorship & Golden Synthesis:** Assigns authoritative source hierarchy (e.g., Salesforce owns billing address, Shopify owns order recency, Zendesk owns satisfaction score).

3. Empirical SME Cost & Time Savings

Traditional Master Data Management implementations (Informatica MDM, Collibra) require 6–9 months of professional services costing between \$150,000 and \$500,000. KaoinAI's automated neural cataloging compresses this deployment cycle to under 48 hours, enabling small data teams to achieve enterprise-grade AI readiness on predictable monthly SaaS budgets.